Multi-Task Learning

Zhiping (Patricia) Xiao

Tuesday Paper-Reading Group

January 21, 2020

UCLA

Overview

1. An Overview of Multi-Task Learning in Deep Neural Networks (ArXiv'17, cited by 521) [1]

2. Fully-adaptive Feature Sharing in Multi-Task Networks with Applications in Person Attribute Classification (CVPR'16)

3. Training Complex Models with Multi-Task Weak Supervision (AAAI'19)

4. Multi-Task Learning Using Uncertainty to Weigh Losses for Scene Geometry and Semantics (CVPR'17)

5. (*) Multi-Task Learning as Multi-Objective Optimization (NeurIPS'18)

---

[1] Originally a blog.

UCLA

Multi-Task Learning (MTL) is often referred to as **joint learning**, **learning with auxiliary tasks**, etc.

It could be applied to either neural models or non-neural models.

The relationship among the tasks matter a lot.

Assume that we have T tasks, each corresponding to a model m_t where $t = 1, 2, \ldots, T$, and d parameters, written as a column vector $a_{.,t}$:

$$a_{.,t} = \begin{bmatrix} a_{1,t} \\ \vdots \\ a_{d,t} \end{bmatrix}$$

Put T columns together we have a matrix $A \in \mathbb{R}^{d \times T}$. a_i is the i^{th} row of A containing the i^{th} feature across all tasks.

Assumption: all models share a small set of features, which means that all but a few **rows** of A are zero.

UCLA

Recall that ℓ_1 norm is also called LASSO (Least Absolute Shrinkage and **Selection Operator**). It is a constraint on the sum of β and thus forces all but a few parameters to be exactly 0. (Explanation comes from <u>here</u>.)

Considering a simple model $y = \beta x$ and $\beta \in \mathbb{R}$, we have a loss function:

$$\mathcal{L} = \min \|y - \beta x\|_2^2$$
$$= \min \left(y^T y - 2\beta y^T x + \|\beta x\|_2^2 \right)$$
$$= \min \left(\beta^2 x^T x - 2\beta y^T x \right)$$

Assuming the least-square solution $\hat{\beta} \geq 0$

To which we could add a norm to regularize β. If it is an ℓ_1 norm we have:

$$\mathcal{L} = \min\left(\beta^2 x^T x - 2\beta y^T x + \lambda|\beta|\right)$$

Taking $\mathcal{L}$'s derivative on β, we have:

$$2x^T x \hat{\beta} - 2y^T x + \lambda sign(\beta) = 0$$

When $\hat{\beta} > 0$ we might be able to reach the optimal point, when it is smaller than zero, since we have an $\hat{\beta} > 0$, we know that $y^T x > 0$, and $\lambda > 0$ by definition; the curve is decreasing (the optimal point $\frac{y^T x + \lambda/2}{x^T x} > 0$ will be on the positive side).

$\therefore$ When using ℓ_1 norm we have:

$$\hat{\beta} = \begin{cases} \frac{y^T x - \lambda/2}{x^T x} & |\beta| > 0, \ y^T x > \lambda/2 \\ 0 & |\beta| \le 0 \end{cases}$$

If it is an ℓ_2 norm instead of ℓ_1, then:

$$2x^T x \hat{\beta} - 2y^T x + \lambda \hat{\beta} = 0$$

$$\hat{\beta} = \frac{y^T x}{x^T x + \lambda/2}$$

Since $\lambda > 0$ and $x^T x \ge 0$, by definition we have $y^T x > 0$, thus $\hat{\beta} > 0$.

UCLA

Revisit the definition of $A \in \mathbb{R}^{d \times T}$, whose t^{th} column is $a_{.,t}$, containing all parameters from task t:

$$a_{.,t} = \begin{bmatrix} a_{1,t} \\ \vdots \\ a_{d,t} \end{bmatrix}$$

i^{th} row of A is a_i, containing the i^{th} feature across all tasks.

Constraints are therefore applied to A to make only a few features non-zero, regularizing by ℓ_1/ℓ_q norm: [2]

$$\left\| \begin{bmatrix} \|a_1\|_q \\ \|a_2\|_q \\ \vdots \\ \|a_d\|_q \end{bmatrix} \right\|_1$$

UCLA

Adding another constraint: [3]

$$\Omega = \|\overline{a}\|^2 + \frac{\lambda}{T} \sum_{t=1}^{T} \|a_{.,t} - \overline{a}\|^2$$

where $\overline{a} = \sum_{t=1}^{T} a_{.,t}/T$.

It enforces the task parameters to be optimized towards their mean, so as to make them more related.

[3]* Learning Multiple Tasks with Kernel Methods (JMLR'05)

UCLA

Figure: Hard parameter sharing. [4]

[4]More popular than soft-parameter sharing.

Figure: Soft parameter sharing.

Architecture

Fully-adaptive Feature Sharing in Multi-Task Networks with Applications in Person Attribute Classification (Yongxi Lu et al.) [5]

- ▶ Base model: CNN.
- ▶ Tasks are pre-defined, the network architecture is automatically-determined.
- ▶ Hard-parameter-sharing.

[5]Person attributes: face, cloth, etc.

UCLA

Figure: Splitting of the layers only applies to one layer at a time, a greedy top-down manner starting from the top layer.

UCLA

How to split the current layer?

How to split the current layer?

By measuring two tasks' (i, j) affinity (similarity) $A(i, j)$, using the votes from **data** samples.

Denote the prediction of the task i for sample n as s_i^n, and the ground-truth label t_i^n, then the error margin is $m_i^n = |t_i^n - s_i^n|$. $m_i \in \mathbb{R}^{\mathcal{N}}$. If for task i, a sample n makes $m_i^n \geq \mathbb{E}\{m_i\}$, we say n is a **difficult example** for task i. Labels are **binary** in this case.

UCLA

How to split the current layer?

By measuring two tasks' (i, j) affinity (similarity) $A(i, j)$, using the votes from **data** samples.

Denote the prediction of the task i for sample n as s_i^n, and the ground-truth label t_i^n, then the error margin is $m_i^n = |t_i^n - s_i^n|$. $m_i \in \mathbb{R}^{\mathcal{N}}$. If for task i, a sample n makes $m_i^n \geq \mathbb{E}\{m_i\}$, we say n is a **difficult example** for task i. Labels are **binary** in this case.

With this definition of possible n's, we define the affinity of two **task** $A(i, j)$ as:

$$A(i, j) = \mathbb{P}(e_i^n = 1, e_j^n = 1) + \mathbb{P}(e_i^n = 0, e_j^n = 0)$$
$$= \mathbb{E}\{e_i^n e_j^n + (1 - e_i^n)(1 - e_j^n)\}$$

UCLA

The affinity of two **branches** (k, l) is $A_b(k, l)$:

$$\tilde{A}_b(k, l) = \text{mean}_{i_k}\left(\min_{j_l} A(i_k, j_l)\right)$$

$$\tilde{A}_b(l, k) = \text{mean}_{i_l}\left(\min_{j_k} A(i_l, j_k)\right)$$

$$A_b(k, l) = \frac{\tilde{A}_b(k, l) + \tilde{A}_b(l, k)}{2}$$

Determined by the largest distance between their associated tasks.

UCLA

There's always a maximum value of clusters c (originally the same as # tasks, will be decreased to the number of clusters in the above layer).

The number of clusters selected is d, $1 \leq d \leq c$. $g_d : [d] \rightarrow [c]$ is the temporary layer's grouping function, associating the current layer with the above one.

The loss function at layer l is (optimized over g_d):

$$L^l(g_d) = (d - 1)L_0 2^{p_l} + \alpha L_S(g_d)$$

where L_0 is the unit cost for branch-creation, p_l is the number of pooling layers above the layer l, larger *branching factor* α encourages more branches, $L_S(g_d)$ is a penalty for separation.

UCLA

For each newly-created branch $i \in [d]$, its separation cost is:

$$L_s^i(g_d) = 1 - \operatorname{mean}_{k \in g^{-1}(i)} \left(\min_{l \in g^{-1}(i)} A_b(k, l) \right)$$

where k, l are branches from the above layer. It measures the maximum distances (minimum affinity) between the tasks within the same group. It penalizes dissimilar tasks grouped into the same branch.

$$L_s(g_d) = \frac{1}{d} \sum_{i \in [d]} L_s^i(g_d)$$
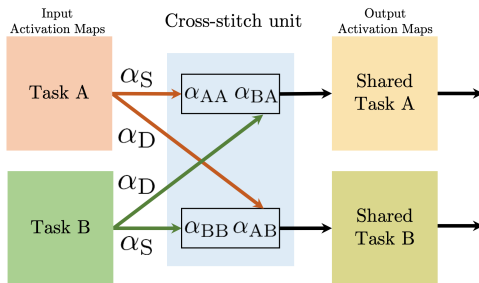
UCLA

Figure: The shared representations is learned by utilizing a linear combination of input activation maps from both tasks, at **each** layer. Cross-stitch units model shared representations as a linear combination of input activation maps.

UCLA

Given two activation maps x_A, x_B from layer l from tasks A and B, we learn **linear** combinations $\tilde{x}_A$, $\tilde{x}_B$ of them parameterized by α, and feed those linear combinations into the next layer. Specifically, at location (i, j):

$$\begin{bmatrix} \tilde{x}_A^{ij} \\ \tilde{x}_B^{ij} \end{bmatrix} = \begin{bmatrix} \alpha_{AA} & \alpha_{AB} \\ \alpha_{BA} & \alpha_{BB} \end{bmatrix} \begin{bmatrix} x_A^{ij} \\ x_B^{ij} \end{bmatrix}$$

The above is called stitching operation.

UCLA

Figure: Applying the above-mentioned cross-stitching to two AlexNet, the units applied only after pooling layers and fully connected layers.

Optimization

Weighted Combination of Loss

Normally MTL's objective is a linear combination of multiple objectives.

Alex Kendall et al. provides a more interpretable estimation by using task-dependent uncertainty [7], estimated by Bayesian modeling, to weigh the losses. We start the computation of $\mathcal{L}(W, \sigma_1, \ldots, \sigma_K)$ from defining multi-task likelihood:

$$p(y|f^W(x)) = \mathcal{N}(y; f^W(x), \sigma^2)$$

$$p(y1, \ldots, y_K|f^W(x)) = p(y_1|f^W(x)) \ldots p(y_K|f^W(x))$$

$$= \mathcal{N}(y_1; f^W(x), \sigma_1^2)\mathcal{N}(y_2; f^W(x), \sigma_2^2)$$

where the distribution $\mathcal{N}(y; f^W(x), \sigma^2)$ is **estimated** by sampling from $Softmax(f^W(x))$.

[7]Claimed to be captured by σ, while data uncertainty by μ.

UCLA

The probability density of observing a single data point x is:

$$p(x; \mu, \sigma) = \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{(x-\mu)^2}{2\sigma^2}\right)$$

The laws of logarithms:

$$\log(ab) = \log(a) + \log(b)$$
$$\log(\frac{a}{b}) = \log(a) - \log(b)$$
$$\log(a^n) = n\log(a)$$

Therefore $(-\log(\sqrt{2\pi})$ is constant and could be removed latter):

$$\log p(x; \mu, \sigma) = -\frac{(x-\mu)^2}{2\sigma^2} - \log(\sigma) - \log(\sqrt{2\pi})$$

UCLA

For a single-task case, the log-likelihood (to be maximized) is:

$$\log p(y|f^W(x)) \propto -\frac{1}{2\sigma^2}\|y - f^W(x)\|^2 - \log \sigma$$

and for multi-task (where there's only 2 tasks), we take the **negative** of the log-likelihood to compute our objective $\mathcal{L}(W, \sigma_1, \sigma_2)$:

$$\mathcal{L}(W, \sigma_1, \sigma_2) \propto \frac{1}{2\sigma_1^2}\|y_1 - f^W(x)\|^2 + \frac{1}{2\sigma_2^2}\|y_2 - f^W(x)\|^2 + \log \sigma_1 \sigma_2$$

Let's define the loss for the two outputs (y_1 and y_2) respectively:

$$\mathcal{L}_1(W) = \|y_1 - f^W(x)\|^2$$
$$\mathcal{L}_2(W) = \|y_2 - f^W(x)\|^2$$

Then we have:

$$\mathcal{L}(W, \sigma_1, \sigma_2) \approx \frac{1}{2\sigma_1^2}\mathcal{L}_1(W) + \frac{1}{2\sigma_2^2}\mathcal{L}_2(W) + \log \sigma_1 \sigma_2$$

Interpret: when σ_i goes up, the relative weight of task i $\lambda_i = \frac{1}{2\sigma_i^2}$ goes down.

(*) There's an extension from Gaussian distribution to Boltzmann (Gibbs) distribution for classification likelihoods' computation in the paper (p 5).

UCLA

Training: In practice, they train the network to predict the **log variance** ($s = \log \sigma^2$), because:

- ▶ the loss avoids any division by zero, thus it is more numerically stable than regressing the variance, σ^2 directly.

- ▶ we can regress unconstrained scalar values, since $\exp(\cdot) > 0$.

UCLA

<u>Ozan Sener & Vladlen Koltun</u> work on solving the problem of *multiple tasks might conflict.*

They used the multiple-gradient-descent algorithm (<u>MGDA</u>) optimizer, together with a definition of the Pareto optimality for MTL (in brief, no other solutions dominants the current solution), they solve the problem by solving the <u>KKT (Karush-Kuhn-Tucker) conditions</u>.

- ▶ Basic Idea: Use multi-objective Karush-Kuhn-Tucker KKT conditions and find a descent direction that decreases all objectives.

- ▶ Can applicable to any problem that uses optimization **based on gradient descent**.

Application

Figure: The goal is to apply a set of weak supervision sources $S = \{s_1, s_2, \ldots, s_m\}$ to an **unlabeled** dataset $\mathcal{X}$ consisting of n data points, and finally apply the label to an end-model $f_w : \mathcal{X} \rightarrow \mathcal{Y}$.

Training Complex Models with Multi-Task Weak Supervision

▶ Feasible Set of input: $\mathcal{X}$, feasible set of labels: $\mathcal{Y}$, *t tasks* in total, *m multi-task weak supervision sources*; X and Y are independent.

▶ $X \in \mathcal{X}$: a data point

▶ $Y \in \mathcal{Y}$: **unobservable** ground-truth task label, $Y = [Y_1, Y_2, \ldots, Y_t]^T$, where for each Y_i, $Y_i \in \{1, 2, \ldots, k_i\}$.

▶ **Examples** of tasks: (1) *Person* v.s. *Organization*; (2) *Doctor* v.s. *Lawyer* (or *N/A*); (3) *Hospital* v.s. *Office* (or *N/A*).

▶ **Tasks' labels** are logically subsumption. e.g. $Y_1 = PERSON, Y_2 = DOCTOR, Y_3 = N/A$ makes a valid $Y \in \mathcal{Y}$. ($Y_2 = DOCTOR$ implies Y_3 and Y_1.)

UCLA

Y is not directly observable, instead, we have access to m multi-task weak supervision sources $s_i \in S$, $i = 1, 2, \ldots, m$; which **emits label vectors** λ_i.

Each λ_i contains labels for some of the tasks, but not for all of them. For the null / abstaining labels, we use 0 to denote. The **coverage set** $\tau_i \subseteq \{1, 2, \ldots, t\}$ is fixed for each s_i, indicating the **tasks** that this source could emit non-zero labels for.

Still consider the previous example, assume that for a certain X the corresponding Y is $[PERSON, DOCTOR, N/A]^T$, we'll probably see $\lambda_i = [PERSON, 0, N/A]^T$, where the coverage set $\tau_i = \{1, 3\}$.

UCLA

There's a non-directional $G_{source} = \{V, E\}$ that is supposed to be given by user who runs this model, where $V = \{Y, \lambda_1, \lambda_2, \ldots, \lambda_m\}$. If (λ_i, λ_j) doesn't exists, it means that the two sources are independent of each other. All (λ_i, Y) always exists.

We define the set of cliques (complete subgraph) of G_{source} as $\mathcal{C}$. Recall that Y is not observable, therefore we define observable cliques set $\mathcal{O} \subset \mathcal{C}$, s.t. we can use it to help with a **matrix completion task**.

(*) The users also need to provide G_{task}, whose structure expresses logical relationships between tasks.

UCLA

Figure: For example, in this case, $\mathcal{C} = \{Y, \lambda_1, \lambda_2, \lambda_3, \lambda_4, \{Y, \lambda_1\}, \{Y, \lambda_2\}, \{Y, \lambda_3\}, \{Y, \lambda_4\}, \{\lambda_1, \lambda_2\}, \{Y, \lambda_1, \lambda_2\}\}$, and $\mathcal{O} = \{\lambda_1, \lambda_2, \lambda_3, \lambda_4, \{\lambda_1, \lambda_2\}\}$

For $C \in \mathcal{C}$, defining $\psi(C)$ as the vector of indicator random variables for **all combinations** of *all but one* of the labels emitted by **each** variable (G_{source} vertices) in clique C, a minimal set of statistics. For $\psi(\mathbf{C})$, it is a collection of all $\psi(C)$ where $C \in \mathbf{C}$.

$$\psi(C, y_C) = \mathbf{1}\{\cap_{i \in C} \; V_i = (y_C)_i | (y_C)_i \in \mathcal{Y}_{\tau_i}\}$$
$$\psi(C) \in \{0, 1\}^{\prod_{i \in C}(|\mathcal{Y}_{\tau_i}| - 1)}$$

The **separator set cliques** of the junction tree (assumed to be unique in a simplification condition) is annotated as $\mathcal{S} \subseteq \mathcal{C}$.

UCLA

What to do next: to learn the **label model** $P_\mu(Y|\lambda)$, which is parameterized by μ, a vector of source correlations and accuracies, and outputs a single probabilistic label vector $\tilde{Y}$.

- $\mu = \mathbb{E}[\psi(\mathcal{C})]$
- $\mathbf{Cov}[\psi(\mathcal{O}), \psi(\mathcal{S})]$ is a function of μ, and it is **unobserved**.

Recall that we said: $\Sigma_{\mathcal{OS}} = \mathbf{Cov}[\psi(\mathcal{O}), \psi(\mathcal{S})]$ is a function of μ, and it is **unobserved**. We also have:

▶ $\Sigma_{\mathcal{O}} = \mathbf{Cov}[\psi(\mathcal{O})]$ be **observed**.

▶ $\Sigma_{\mathcal{S}} = \mathbf{Cov}[\psi(\mathcal{S})] = \mathbf{Cov}[\psi(Y)]$ is **estimable**, it is a function of $P(Y)$.

▶ The mathematical calculation is based on:

$$\mathbf{Cov}[\psi(\mathcal{O} \cup \mathcal{S})] \equiv \Sigma = \begin{bmatrix} \Sigma_{\mathcal{O}} & \Sigma_{\mathcal{OS}} \\ \Sigma_{\mathcal{OS}}^{T} & \Sigma_{\mathcal{S}} \end{bmatrix}$$

All the remaining is matrix computation details, working on matrix completion. For details, please visit the longer version of this paper on ArXiv.

UCLA